

WILLIAM T. FARRINGTON, M.D.

williamtfarrington.com — william.t.farrington@gmail.com

I am a board-certified internist and acute-care hospitalist driven by a well-known engineering maxim: “all models are wrong, but some are useful.” After studying engineering and medicine at Texas A&M and completing my residency in Dallas, I spent two years on the McGovern Medical School faculty at Memorial Hermann Hospital–Texas Medical Center. I concluded this appointment in spring of 2026 to relocate to Philadelphia with my fiancée. Now licensed in Pennsylvania, I continue public and private data work while seeking a clinical home in which to combine inpatient care, teaching, and responsible technical work.

At Memorial Hermann-TMC, I practiced high-volume, high-acuity medicine in a hospitalist group responsible for more than 40 percent of admissions. My primary assignments spanned oncology, neurology, transplant, geriatrics, and trauma service lines. Fulfilling my group’s triage and capacity-management duties involved ED- and ICU-facing acuity forecasting and level-of-care decision-making, as well as handling transfer requests from outside institutions. I was met continually by the collective complexities of inpatient medicine: mediating discussions regarding primary and consultant roles, balancing attending-level communication with opportunities for trainee education, and negotiating for responsible use of resources. But complexity was not always a curse; when prospective clinical heuristics, inter-departmental protocols, and system-level models reconciled into a “willful architecture,” the clinical overwhelm would recede for a time. This wholly romantic phenomenon informs my most ambitious development project, *mimicwarehouse*: a local data lab built over the MIMIC-IV corpus to study deterioration, peri-critical care, and resource use.

My fascination with system-level thinking also permeated my teaching efforts. I did my best while supervising admissions and other resident work to evaluate through discrete lenses: acute risk mitigation, multidisciplinary team coordination, and coherence of the management plan. But periodically the “most constructive” feedback felt interstitial, as if it existed in the space between otherwise acceptable inferences and actions. I addressed these gaps in classical formats, delivering annual triage lectures for incoming interns, disease-management lectures for surgical residents, and the medication reconciliation aspect of our “safe discharge” workshops. I also realized that teaching a modern approach to patient care meant building system models and interfaces which comprise the patient, the hospital, the city, the electronic health record, and an ever-expanding landscape of internal and open-access AI tools. In these last two domains there is still much authority to be claimed by attending physicians, particularly regarding what “clinically and technologically responsible use” looks like. Attending physicians must adopt these tools early and earnestly enough to catalog and characterize the hazards, to lobby for tools that facilitate privacy and quality before profit, and to widen the educational scope. In this spirit, I am actively developing simulation-based projects like *medrecsim*, a medication reconciliation interface and a resident-facing example of what physician-driven software design can produce.

My ongoing engineering work leaves me unwilling to eschew AI tools or accept them casually. Pilot deployments of Viz.ai, Ambiance, and Regard models for the electronic health record at my prior institution only reinforced *de facto* wisdom: these tools play a very small part in the full-scale operation of a safe and efficient inpatient system. They also represent poorly differentiated and under-acknowledged “risk surfaces.” I choose to define risk in my work by what software may read and what its output may induce. But clinically facing tools must also undergo design ratification, iterative validation, and accountable review. Of all my public projects, only *mimicwarehouse* faces de-identified patient data, with records maintained locally and redundantly inaccessible to “agentic context.” Other clinically adjacent prototypes use synthetic or fictional inputs. I bring strong engineering fundamentals to my work, but I could never downplay the role of these tools in my private data work or the innate caution their use provokes in me. I view them as potentially dangerous but ultimately democratizing forces staged to empower the pure academicians, practitioners, and budding “physicians of tomorrow” who are most intimately familiar with their under-modeled “home systems.”

Physicians are united by a sovereign dedication to quality patient care. This should never change. Even as I strive for greater responsibility in operational informatics and collaborative data work, I am confident that my practice could never leave the bedside. I am guided by the late Paul Farmer’s mantric reminder to “fight the long defeat,” and the most important battles, winnable or not, always seem to be fought away from the computer screen. I feel compelled to play my small part in shaping the future generation of physicians, to serve as an example of evidence-based practice as often as I can and refract my insights through the engineering prism when appropriate. I am neither a tenured professor nor a polished informatician, but I love teaching and “data wrangling” almost as much as I love the messy groundwork of hospital medicine.